

# CSI 62

## Operating Systems and Systems Programming

### Lecture 11

## Scheduling (finished), Deadlock

March 3<sup>rd</sup>, 2020

Prof. John Kubiatawicz

<http://cs162.eecs.Berkeley.edu>

*Acknowledgments: Lecture slides are from the Operating Systems course taught by John Kubiatawicz at Berkeley, with few minor updates/changes. When slides are obtained from other sources, a reference will be noted on the bottom of that slide, in which case a full list of references is provided on the last slide.*

# Recall: Scheduling Policy Goals/Criteria

---

- Minimize Response Time
  - Minimize elapsed time to do an operation (or job)
  - Response time is what the user sees:
    - » Time to echo a keystroke in editor
    - » Time to compile a program
    - » Real-time Tasks: Must meet deadlines imposed by World
- Maximize Throughput
  - Maximize operations (or jobs) per second
  - Throughput related to response time, but not identical:
    - » Minimizing response time will lead to more context switching than if you only maximized throughput
  - Two parts to maximizing throughput
    - » Minimize overhead (for example, context-switching)
    - » Efficient use of resources (CPU, disk, memory, etc)
- Fairness
  - Share CPU among users in some equitable way
  - Fairness is not minimizing average response time:
    - » Better *average* response time by making system *less* fair


# Recall: What if we Knew the Future?

---

- Could we always mirror best FCFS?
- Shortest Job First (SJF):
  - Run whatever job has the least amount of computation to do
  - Sometimes called “Shortest Time to Completion First” (STCF)
- Shortest Remaining Time First (SRTF):
  - Preemptive version of SJF: if job arrives and has a shorter time to completion than the remaining time on the current job, immediately preempt CPU
  - Sometimes called “Shortest Remaining Time to Completion First” (SRTCF)
- These can be applied either to a whole program or the current CPU burst of each program
  - Idea is to get short jobs out of the system
  - Big effect on short jobs, only small effect on long ones
  - Result is better average response time



# Recall: Multi-Level Feedback Scheduling



- Another method for exploiting past behavior
  - First used in CTSS
  - Multiple queues, each with different priority
    - » Higher priority queues often considered “foreground” tasks
  - Each queue has its own scheduling algorithm
    - » e.g. foreground – RR, background – FCFS
    - » Sometimes multiple RR priorities with quantum increasing exponentially (highest: 1ms, next: 2ms, next: 4ms, etc)
- Adjust each job’s priority as follows (details vary)
  - Job starts in highest priority queue
  - If timeout expires, drop one level
  - If timeout doesn’t expire, push up one level (or to top)


# Real-Time Scheduling (RTS)

---

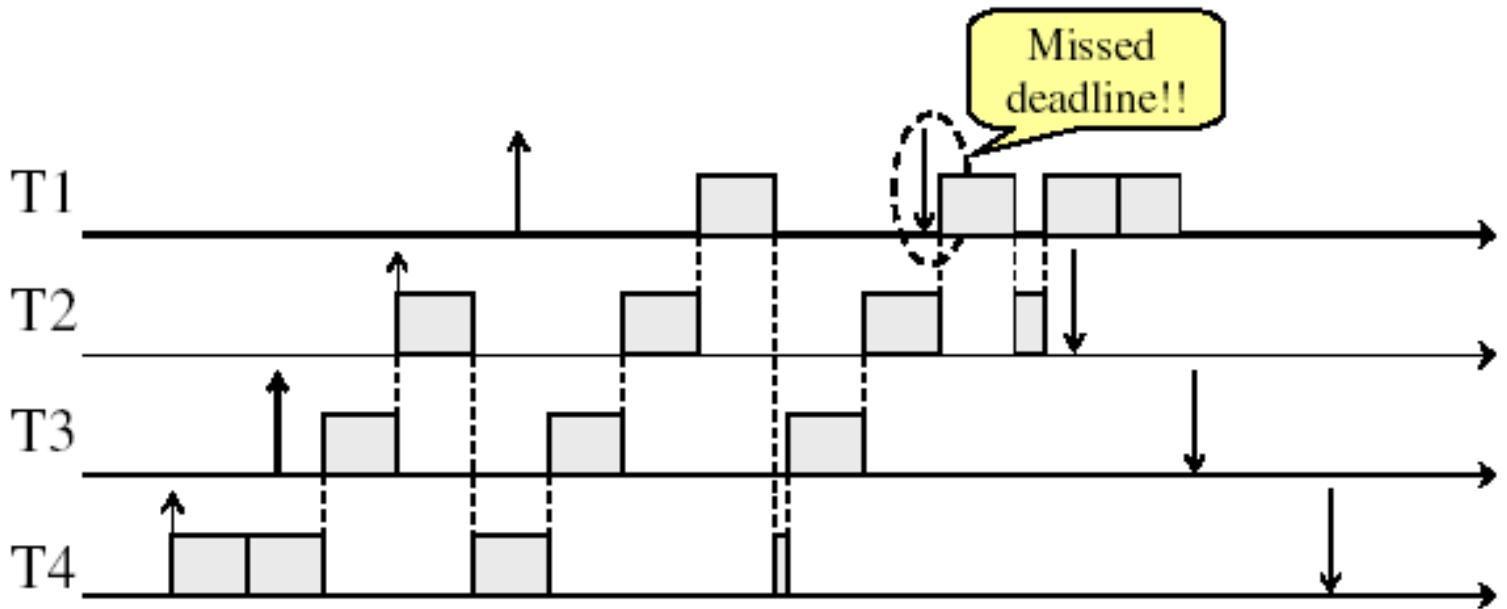
- Efficiency is important but **predictability** is essential:
  - We need to predict with confidence worst case response times for systems
  - In RTS, performance guarantees are:
    - » Task- and/or class centric and often ensured a priori
  - In conventional systems, performance is:
    - » System/throughput oriented with post-processing (... wait and see ...)
  - Real-time is about enforcing predictability, and does not equal fast computing!!!
- Hard Real-Time
  - *Attempt to meet all deadlines*
  - EDF (Earliest Deadline First), LLF (Least Laxity First), RMS (Rate-Monotonic Scheduling), DM (Deadline Monotonic Scheduling)
- Soft Real-Time
  - *Attempt to meet deadlines with high probability*
  - Minimize miss ratio / maximize completion ratio (firm real-time)
  - Important for multimedia applications
  - CBS (Constant Bandwidth Server)

# Recall: Realtime Workload Characteristics

- Tasks are preemptable, independent with arbitrary arrival (=release) times
- Tasks have deadlines ( $D$ ) and known computation times ( $C$ )
- Example Setup:




\_\_\_\_\_



# Recall: Earliest Deadline First (EDF)

- Tasks **periodic** with period  $P$  and computation  $C$  in each period:  $(P_i, C_i)$  for each task  $i$
- Preemptive priority-based dynamic scheduling:
  - Each task is assigned a (current) priority based on how close the absolute deadline is (i.e.  $D_i^{t+1} = D_i^t + P_i$  for each task!)
  - The scheduler always schedules the active task with the closest absolute deadline



- Schedulable when  $\sum_{i=1}^n \left( \frac{C_i}{P_i} \right) \leq 1$




# Choosing the Right Scheduler

---


I Care About:	Then Choose:
CPU Throughput	FCFS
Avg. Response Time	SRTF Approximation
I/O Throughput	SRTF Approximation
Fairness (CPU Time)	Linux CFS
Fairness - Wait Time to Get CPU	Round Robin
Meeting Deadlines	EDF
Favoring Important Tasks	Priority

# A Final Word On Scheduling


- When do the details of the scheduling policy and fairness really matter?
  - When there aren't enough resources to go around
- When should you simply buy a faster computer?
  - (Or network link, or expanded highway, or ...)
  - One approach: Buy it when it will pay for itself in improved response time
    - » Perhaps you're paying for worse response time in reduced productivity, customer angst, etc...
    - » Might think that you should buy a faster X when X is utilized 100%, but usually, response time goes to infinity as utilization  $\Rightarrow$  100%
- An interesting implication of this curve:
  - Most scheduling algorithms work fine in the “linear” portion of the load curve, fail otherwise
  - Argues for buying a faster X when hit “knee” of curve



# Starvation vs Deadlock



- **Starvation**: thread waits indefinitely
  - Example, **low**-priority thread waiting for resources constantly in use by high-priority threads
- **Deadlock**: circular **waiting** for resources
  - Thread A owns Res 1 and is waiting for Res 2
  - Thread B owns Res 2 and is waiting for Res 1



- Deadlock  $\Rightarrow$  Starvation but not **vice versa**
  - Starvation can end (but doesn't have to)
  - Deadlock can't end without external intervention


# Example: Single-Lane Bridge Crossing

---



**CA 140 to Yosemite National Park**

# Bridge Crossing Example



- Each segment of road can be **viewed** as a **resource**
  - Car must own the segment under them
  - Must acquire segment that they are moving into
- For bridge: must acquire both halves
  - Traffic only in one direction at a time
  - Problem occurs when two cars in opposite directions on bridge: each acquires one segment and needs next
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
  - Several cars may have to be backed up
- Starvation is possible
  - East-going traffic really fast  $\Rightarrow$  no one goes west

# One Lane Bridge Revisited: Deadlock with Locks

---


## Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```

## Thread B

```
y.Acquire();  
x.Acquire();  
...  
x.Release();  
y.Release();
```

## Nondeterministic Deadlock



# Deadlock with Locks: Unlucky Case

## Thread A

**x.Acquire();**

**y.Acquire();**

**<stalled>**

**<unreachable>**

...

**y.Release();**

**x.Release();**

## Thread B

**y.Acquire();**


**x.Acquire(); <stalled>**

**<unreachable>**

...

**x.Release();**

**y.Release();**



# Deadlock with Locks: “Lucky” Case

---

## Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```

## Thread B


```
y.Acquire();  
  
x.Acquire();  
...  
x.Release();  
y.Release();
```

**Sometimes schedule won't trigger deadlock**



# Train Example (Wormhole-Routed Network)

- Circular dependency (Deadlock!)
  - Each train wants to turn right
  - Blocked by other trains
  - Similar problem to multiprocessor networks
- Fix? Imagine grid extends in all four directions
  - Force ordering of channels (tracks)
    - » Protocol: Always go east-west first, then north-south
  - Called “dimension ordering” (X then Y)



# Other Types of Deadlock

---

- Threads often **block** waiting for **resources**
  - Locks
  - Terminals
  - Printers
  - CD drives
  - Memory
- Threads often block waiting for other threads
  - Pipes
  - Sockets
- You can deadlock on any of these!

# Deadlock with Space

---

## Thread A

**AllocateOrWait(1 MB)**

**AllocateOrWait(1 MB)**

**Free(1 MB)**

**Free(1 MB)**

## Thread B

**AllocateOrWait(1 MB)**


**AllocateOrWait(1 MB)**

**Free(1 MB)**

**Free(1 MB)**

If only 2 MB of space, we get same deadlock situation

# Dining Lawyers Problem



- Five chopsticks/Five lawyers (really cheap restaurant)
  - Free-for all: Lawyer will grab any one they can
  - Need two chopsticks to eat
- What if all grab at same time?
  - Deadlock!
- How to fix deadlock?
  - Make one of them give up a chopstick (Hah!)
  - Eventually everyone will get chance to eat
- How to prevent deadlock?
  - Never let lawyer take last chopstick if no hungry lawyer has two chopsticks afterwards

# Four requirements for occurrence of Deadlock


---

- Mutual exclusion
  - Only one thread at a time can use a resource.
- Hold and wait
  - Thread holding at least one resource is waiting to acquire additional resources held by other threads
- No preemption
  - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
- Circular wait
  - There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
    - »  $T_1$  is waiting for a resource that is held by  $T_2$
    - »  $T_2$  is waiting for a resource that is held by  $T_3$
    - » ...
    - »  $T_n$  is waiting for a resource that is held by  $T_1$

# Detecting Deadlock: Resource-Allocation Graph

- System Model

- A set of Threads  $T_1, T_2, \dots, T_n$
- Resource types  $R_1, R_2, \dots, R_m$ 
  - CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances
- Each thread utilizes a resource as follows:
  - » Request () / Use () / Release ()




- Resource-Allocation Graph:


- $V$  is partitioned into two types:
  - »  $T = \{T_1, T_2, \dots, T_n\}$ , the set threads in the system.
  - »  $R = \{R_1, R_2, \dots, R_m\}$ , the set of resource types in system
- request edge – directed edge  $T_i \rightarrow R_j$
- assignment edge – directed edge  $R_j \rightarrow T_i$

# Resource-Allocation Graph Examples


- Model:
  - request edge – directed edge  $T_i \rightarrow R_j$
  - assignment edge – directed edge  $R_i \rightarrow T_j$



Simple Resource  
Allocation Graph




Allocation Graph  
With Deadlock



Allocation Graph  
With Cycle, but  
No Deadlock

# Deadlock Detection Algorithm

- Only one of each type of resource  $\Rightarrow$  look for loops
  - More General Deadlock Detection Algorithm
    - Let  $[X]$  represent an m-array vector of non-negative integers (quantities of resources of each type):
      - [FreeResources]:** Current free resources each type
      - [Request<sub>x</sub>]:** Current requests from thread X
      - [Alloc<sub>x</sub>]:** Current resources held by thread X
    - See if tasks can eventually terminate on their own
- ```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```
- Nodes left in **UNFINISHED**  $\Rightarrow$  deadlocked





# How should a system deal with deadlock?

---

- Four different approaches:
  1. Deadlock prevention: write your code in a way that it **isn't prone** to deadlock
  2. Deadlock recovery: let deadlock **happen**, and then figure out how to **recover** from it
  3. Deadlock avoidance: dynamically **delay** resource requests so deadlock **doesn't** happen
  4. Deadlock denial: **ignore** the possibility of deadlock
- Modern operating systems:
  - Make sure the system **isn't involved** in any deadlock
  - **Ignore deadlock** in applications
    - » “**Ostrich** Algorithm”

# Techniques for Preventing Deadlock

---

- Infinite resources
  - Include enough resources so that no one ever runs out of resources. Doesn't have to be infinite, just large
  - Give illusion of infinite resources (e.g. virtual memory)
  - Examples:
    - » Bay bridge with 12,000 lanes. Never wait!
    - » Infinite disk space (not realistic yet?)
- No Sharing of resources (totally independent threads)
  - Not very realistic
- Don't allow waiting
  - How the phone company avoids deadlock
    - » Call to your Mom in Toledo, works its way through the phone lines, but if blocked get busy signal.
  - Technique used in Ethernet/some multiprocessor nets
    - » Everyone speaks at once. On collision, back off and retry
  - Inefficient, since have to keep retrying
    - » Consider: driving to San Francisco; when hit traffic jam, suddenly you're transported back home and told to retry!

# (Virtually) Infinite Resources

---

## Thread A

**AllocateOrWait(1 MB)**

**AllocateOrWait(1 MB)**

**Free(1 MB)**

**Free(1 MB)**

## Thread B

**AllocateOrWait(1 MB)**

**AllocateOrWait(1 MB)**

**Free(1 MB)**

**Free(1 MB)**

**With virtual memory we have “infinite” space so everything will just succeed.**

# Techniques for Preventing Deadlock

---

- Make all threads request everything they'll need at the beginning.
  - Problem: Predicting future is hard, tend to over-estimate resources
  - Example:
    - » If need 2 chopsticks, request both at same time
    - » Don't leave home until we know no one is using any intersection between here and where you want to go; only one car on the Bay Bridge at a time
- Force all threads to request resources in a particular order preventing any cyclic use of resources
  - Thus, preventing deadlock
  - Example (`x.Acquire()`, `y.Acquire()`, `z.Acquire()`,...)
    - » Make tasks request disk, then memory, then...
    - » Keep from deadlock on freeways around SF by requiring everyone to go clockwise

# Request Resources Atomically (I)

---

Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```

Thread B

```
y.Acquire();  
x.Acquire();  
...  
x.Release();  
y.Release();
```

Consider instead:

Thread A

```
Acquire_both(x, y);  
...  
y.Release();  
x.Release();
```

Thread B

```
Acquire_both(y, x);  
...  
x.Release();  
y.Release();
```

## Request Resources Atomically (2)

---

Or consider this:

Thread A

**z.Acquire();**

**x.Acquire();**

**y.Acquire();**

**z.Release();**

...

**y.Release();**

**x.Release();**

Thread B

**z.Acquire();**

**y.Acquire();**

**x.Acquire();**

**z.Release();**

...

**x.Release();**

**y.Release();**

# Acquire Resources in Consistent Order

---

## Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```

## Thread B

```
y.Acquire();  
x.Acquire();  
...  
x.Release();  
y.Release();
```

## Consider instead:

## Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```


## Thread B

```
x.Acquire();  
y.Acquire();  
...  
x.Release();  
y.Release();
```

Does it matter in  
which order the  
locks are released?

# Review: Train Example (Wormhole-Routed Network)

- Circular dependency (Deadlock!)
  - Each train wants to turn right
  - Blocked by other trains
  - Similar problem to multiprocessor networks
- Fix? Imagine grid extends in all four directions
  - Force ordering of channels (tracks)
    - » Protocol: Always go east-west first, then north-south
  - Called “dimension ordering” (X then Y)





# Techniques for Recovering from Deadlock

---

- **Terminate** thread, force it to give up resources
  - In Bridge example, Godzilla picks up a car, hurls it into the river. Deadlock solved!
  - Hold dining lawyer in contempt and take away in handcuffs
  - But, not always possible – killing a thread holding a mutex leaves world inconsistent
- **Preempt** resources without killing off thread
  - **Take away** resources from thread temporarily
  - Doesn't always fit with semantics of computation
- **Roll back** actions of deadlocked threads
  - Hit the **rewind** button on TiVo, **pretend** last few minutes never happened
  - For bridge example, make one car **roll backwards** (may require others behind him)
  - Common technique in databases (**transactions**)
  - Of course, if you **restart** in exactly the same way, may reenter deadlock once again
- Many operating systems use other options

# Pre-empting Resources

---

## Thread A

**AllocateOrWait(1 MB)**

**AllocateOrWait(1 MB)**

**Free(1 MB)**

**Free(1 MB)**

## Thread B

**AllocateOrWait(1 MB)**

**AllocateOrWait(1 MB)**

**Free(1 MB)**

**Free(1 MB)**

**With virtual memory we have “infinite” space so everything will just succeed.**

**Alternative view: we are “pre-empting” memory when paging out to disk, and giving it back when paging back in**

# Techniques for Deadlock Avoidance

---

- Idea: When a thread requests a resource, OS checks if it would result in deadlock
  - If not, it grants the resource right away
  - If so, it waits for other threads to release resources

**THIS DOES NOT WORK!!!!**

- Example:

## Thread A

Blocks... **x.Acquire();**  
**y.Acquire();**  
...  
**y.Release();**  
**x.Release();**

## Thread B

**y.Acquire();** Wait...  
**x.Acquire();** But it's too late...  
...  
**x.Release();**  
**y.Release();**

# Deadlock Avoidance: Three States

---

- Safe state
  - System can delay resource acquisition to prevent deadlock
- Unsafe state
  - No deadlock yet...
  - But threads can request resources in a pattern that unavoidably leads to deadlock
- Deadlocked state
  - There exists a deadlock in the system
  - Also considered “unsafe”

**Deadlock avoidance:  
prevent system from  
reaching an unsafe state**

# Deadlock Avoidance

---

- Idea: When a thread requests a resource, OS checks if it would result in ~~deadlock~~ **an unsafe state**
  - If not, it grants the resource right away
  - If so, it waits for other threads to release resources
- Example:

## Thread A

```
x.Acquire();  
y.Acquire();  
...  
y.Release();  
x.Release();
```

## Thread B

```
y.Acquire();  
x.Acquire();  
...  
x.Release();  
y.Release();
```

Wait until  
Thread A  
releases the  
mutex

# Banker's Algorithm for Avoiding Deadlock

- Toward right idea:
  - State maximum (max) resource needs in advance
  - Allow particular thread to proceed if:  
 $(\text{available resources} - \text{\#requested}) \geq \text{max}$   
remaining that might be needed by any thread
- Banker's algorithm (less conservative):
  - Allocate resources dynamically
    - » Evaluate each request and grant if some ordering of threads is still deadlock free afterward
    - » Technique: pretend each request is granted, then run deadlock detection algorithm, substituting  
 $([\text{Max}_{\text{node}}] - [\text{Alloc}_{\text{node}}] \leq [\text{Avail}])$  for  $([\text{Request}_{\text{node}}] \leq [\text{Avail}])$   
Grant request if result is deadlock free (conservative!)



# Banker's Algorithm for Avoiding Deadlock

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

// technique: pretend each request is granted, then run deadlock detection algorithm, substituting  
([Max<sub>node</sub>] - [Alloc<sub>node</sub>] <= [Avail]) for ([Request<sub>node</sub>] <= [Avail])  
Grant request if result is deadlock free (conservative!)



# Banker's Algorithm for Avoiding Deadlock

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ( [Maxnode] - [Allocnode] <= [Avail] ) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

// technique: pretend each request is granted, then run deadlock detection algorithm, substituting  
([Max<sub>node</sub>] - [Alloc<sub>node</sub>] <= [Avail]) for ([Request<sub>node</sub>] <= [Avail])  
Grant request if result is deadlock free (conservative!)






# Banker's Algorithm for Avoiding Deadlock

- Toward right idea:
  - State maximum resource needs in advance
  - Allow particular thread to proceed if:  
 $(\text{available resources} - \# \text{requested}) \geq \text{max remaining that might be needed by any thread}$
- Banker's algorithm (less conservative):
  - Allocate resources dynamically
    - » Evaluate each request and grant if some ordering of threads is still deadlock free afterward
    - » Technique: pretend each request is granted, then run deadlock detection algorithm, substituting  
 $([\text{Max}_{\text{node}}] - [\text{Alloc}_{\text{node}}] \leq [\text{Avail}])$  for  $([\text{Request}_{\text{node}}] \leq [\text{Avail}])$   
Grant request if result is deadlock free (conservative!)
    - » Keeps system in a "SAFE" state, i.e. there exists a sequence  $\{T_1, T_2, \dots, T_n\}$  with  $T_1$  requesting all remaining resources, finishing, then  $T_2$  requesting all remaining resources, etc..
  - Algorithm allows the sum of maximum resource needs of all current threads to be greater than total resources



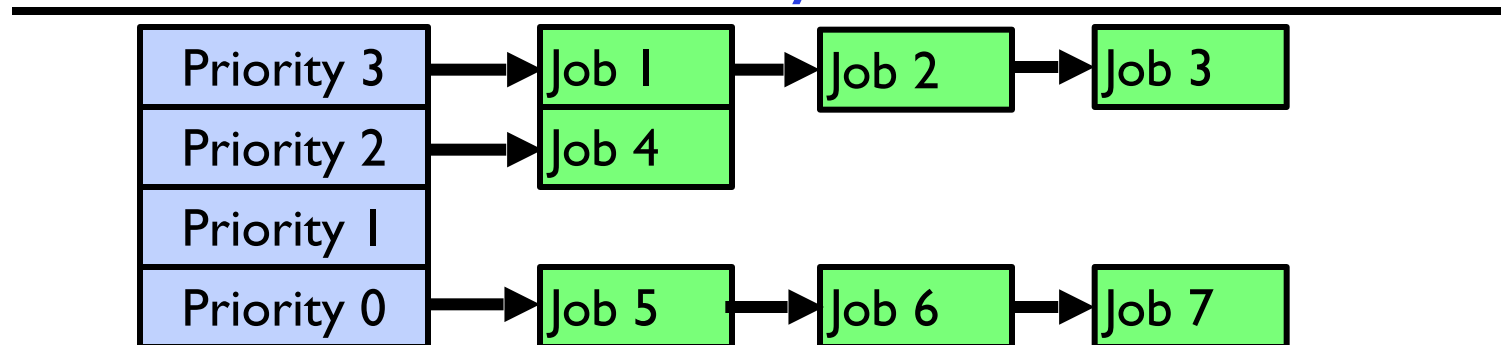
# Banker's Algorithm Example



- Banker's algorithm with dining lawyers
  - “Safe” (won't cause deadlock) if when try to grab chopstick either:
    - » Not last chopstick
    - » Is last chopstick but someone will have two afterwards
  - What if k-handed lawyers? Don't allow if:
    - » It's the last one, no one would have k
    - » It's 2<sup>nd</sup> to last, and no one would have k-1
    - » It's 3<sup>rd</sup> to last, and no one would have k-2
    - » ...



## Recall: Priority Scheduler



- Execution Plan
  - Always execute highest-priority runnable jobs to completion
  - Each queue can be processed in RR with some time-quantum
- Problems:
  - Starvation:
    - » Lower priority jobs don't get to run because of higher priority jobs
  - Priority Inversion:
    - » Not strictly a problem with priority scheduling, but happens when low priority task has lock needed by high-priority task
    - » Usually involves third, intermediate priority task that keeps running even though high-priority task should be running
  - Are either of these problems examples of DEADLOCK?

# Priority Donation as a remedy to Priority Inversion

- Does Priority Inversion cause Deadlock? Not usually.
- Consider:
  - 3 threads, T1, T2, T3 in priority order (T3 highest)
  - T1 grabs lock, T3 tries to acquire, then sleeps, T2 running
  - Will this make progress?
    - » No, as long as T2 is running
    - » But T2 could stop at any time and the problem would resolve itself...
    - » So, this is *not* a deadlock (it is a livelock). But it could last a long time...
  - Why is this a priority inversion?
    - » T3 is prevented from running by T2
- What is priority donation?
  - When high priority Thread TB is about to sleep while waiting for a lock held by lower priority Thread TA, it may temporarily donate its priority to the holder of the lock if that lock holder has a lower priority
    - » So,  $\text{Priority}(\text{TB}) \Rightarrow \text{TA}$  until lock is released
  - So, now, TA runs with high priority until it releases its lock, at which time its priority is restored to its original priority.
- How does priority donation help both above priority inversion scenario?
  - Briefly raising T1 to the same priority as T3  $\Rightarrow$  T1 can run and release lock, allowing T3 to run
  - Does priority donation involve taking lock away from T1?
    - » NO! That would break semantics of the lock and potentially corrupt any information protected by lock!

# Summary

---

- Real-time scheduling
  - Need to meet a deadline, predictability essential
  - Earliest Deadline First (EDF) and Rate Monotonic (RM) scheduling
- Starvation vs. Deadlock
  - Starvation: thread waits indefinitely
  - Deadlock: circular waiting for resources
- Four conditions for deadlocks
  - Mutual exclusion
  - Hold and wait
  - No preemption
  - Circular wait
- Techniques for addressing Deadlock
  - Deadlock prevention:
    - » write your code in a way that it isn't prone to deadlock
  - Deadlock recovery:
    - » let deadlock happen, and then figure out how to recover from it
  - Deadlock avoidance:
    - » dynamically delay resource requests so deadlock doesn't happen
    - » Banker's Algorithm provides an algorithmic way to do this
  - Deadlock denial:
    - » ignore the possibility of deadlock